

# COMPSCI 589

## Lecture 5: Support Vector Machines, Basis Expansion and Kernels

Benjamin M. Marlin

College of Information and Computer Sciences  
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).  
Created with support from National Science Foundation Award# IIS-1350522.

# Overview

- To date we've seen one example of a discriminative linear classifier.
- Today we'll introduce a second example, support vector machines.
- We'll then address the question of how to increase the capacity of linear classifiers so they can produce non-linear classification boundaries.

# Support Vector Machines

- A binary support vector machine is a discriminative classifier that takes labels in the set  $\{-1, 1\}$ .
- The decision function has the form:

$$f_{SVM}(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x} + b)$$

- It's easy to show that the decision boundary for logistic regression can be written in exactly the same way.
- **Question:** If logistic regression and SVMs have the same form for their decision boundaries, how do they differ?

# Logistic Loss

- In the case of logistic regression with  $\ell_2$  regularization, we select the model parameters by maximizing the function:

$$C \sum_{i=1}^n \log P(Y = y_i | \mathbf{X} = \mathbf{x}_i) - \|\mathbf{w}\|_2^2$$

- Under the assumption that the labels take the values  $\{-1, 1\}$ , it can be shown that this is equivalent to minimizing the function:

$$C \sum_{i=1}^n \log(1 + \exp(-y_i \cdot g(\mathbf{x}))) + \|\mathbf{w}\|_2^2$$

where  $L_{\log}(y_i, g(\mathbf{x}_i)) = \log(1 + \exp(-y_i \cdot g(\mathbf{x})))$  is the *logistic loss function* and  $g(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$ .

# Hinge Loss

- In the case of SVMs with  $\ell_2$  regularization, we select the model parameters by minimizing the function:

$$C \sum_{i=1}^n \max(0, 1 - y_i \cdot g(\mathbf{x}_i)) + \|\mathbf{w}\|_2^2$$

- The function  $L_h(y_i, g(\mathbf{x}_i)) = \max(0, 1 - y_i \cdot g(\mathbf{x}_i))$  is called the *hinge loss*.

# Zero-One Loss



- Both the logistic loss and the hinge loss are convex upper bounds on the zero-one loss:

$$L_{01}(y_i, g(\mathbf{x}_i)) = \mathbb{I}[y_i \neq \text{sign}(g(\mathbf{x}_i))]$$

- The average zero-one loss over a data set is exactly the classification error rate.
- This is the loss function we'd like to minimize, but this generally isn't computationally feasible, thus the need for surrogate loss functions.
- Hinge loss has some advantages over logistic loss, as we'll see.


# Maximum Margin Property

Part of popularity of SVMs stems from the fact that the hinge loss results in the *maximum margin* decision boundary when the training cases are linearly separable.



# Support Vector Property

In the linearly separable case, some data points will always fall exactly on the margins. These points are called *support vectors*.





# SVMs vs Logistic Regression

- SVMs and Logistic regression are both discriminative linear classifiers with identical capacity and space complexity.
- SVMs and Logistic regression have very similar convex loss functions and identical regularizers.
- The hinge loss is not differentiable, unlike the logistic loss, so SVMs require more advanced optimization methods (sub-gradient descent or quadratic programming), but there are extremely good algorithms and implementations available.
- The maximum-margin property of SVMs can yield better generalization than logistic regression when data is scarce.

# SVMs vs Logistic Regression

- It is somewhat more difficult to form a multi-class SVM than a multi-class logistic regression model, and many implementations use “hacks” like one-vs-all.
- Unlike logistic regression, SVMs do not produce probabilistic outputs, but again, there are hacks that can estimate probabilities.

# The Problem with Linear Models

- The problem with linear classifiers is that their decision boundaries are by definition linear in the input feature space.
- This means that they can have very high bias on complex data.
- **Question:** How can we relax the constraint of linear decision boundaries while retaining the nice properties of linear classifiers?

# Basis Function Expansion

- One very simple solution is to apply a set of functions  $\phi_1, \dots, \phi_K$  to the raw feature vector  $\mathbf{x}$  to map it in to a new feature space:

$$\phi(\mathbf{x}) = [\phi_1(\mathbf{x}), \dots, \phi_K(\mathbf{x})]$$

- This is called a *basis function expansion* since  $K > D$  in general. This requires that we know the functions  $\phi_1, \dots, \phi_K$  that we want to apply in advance.
- We then define a linear classifier (SVM or Logistic Regression) in this new feature space:

$$\mathbf{w}^T \phi(\mathbf{x}) + b$$

# Basis Function Expansion Examples

- **Degree 2 Polynomial Basis:** We include all single features  $x_d$ , their squares  $x_d^2$ , and all products of two distinct features  $x_d x_{d'}$ .
- **Degree  $B$  Polynomial Basis:** We include all single features  $x_d$ , and all unique products of between 2 and  $B$  features.
- The problem is that the space complexity of representing the expanded set of features is essentially  $O(D^B)$ .
- Next we'll see how this problem can be solved.

# Representer Theorem

- One of the interesting properties of SVMs is that the optimal weight vectors can always be expressed as a weighted linear combination of the data vectors:

$$\mathbf{w} = \sum_{j=1}^N \alpha_j \mathbf{x}_j$$

- This result is called the *representer theorem*.

# Dependence on Inner Products

- Plugging this result back in to the SVM objective we find that the objective only depends on the data through inner products:  $\mathbf{x}_j^T \mathbf{x}_i$ :

$$g(\mathbf{x}_i) = \mathbf{w}^T \mathbf{x}_i + b = \sum_{j=1}^N \alpha_j \mathbf{x}_j^T \mathbf{x}_i + b$$

$$\|\mathbf{w}\|_2^2 = \mathbf{w}^T \mathbf{w} = \sum_{j=1}^N \sum_{i=1}^N \alpha_j \alpha_i \mathbf{x}_j^T \mathbf{x}_i$$

# Basis Expansion and Representer Theorem

- Under an arbitrary basis expansion  $\phi(\mathbf{x})$  this result becomes:

$$g(\phi(\mathbf{x}_i)) = \sum_{j=1}^N \alpha_j \phi(\mathbf{x}_j)^T \phi(\mathbf{x}_i) + b$$
$$\|\mathbf{w}\|_2^2 = \sum_{j=1}^N \sum_{i=1}^N \alpha_j \alpha_i \phi(\mathbf{x}_j)^T \phi(\mathbf{x}_i)$$

- It can be shown in the linearly separable case that the  $\alpha_i$  parameters for data cases that are not support vectors are always 0.



# The Kernel Trick

- Amazingly, for many useful basis function expansions  $\phi(\mathbf{x})$ , it is possible to find a function  $\mathcal{K}(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})^T \phi(\mathbf{x}')$  that can compute the inner product under the basis expansion **without ever explicitly performing the basis function expansion!**
- Such functions are called *kernel functions* and this is known as the “Kernel Trick”.
- Importantly, you can also directly learn the parameters  $\alpha_i$  and  $b$  using the kernel trick, without constructing the basis expansion.
- Interestingly, there exist kernels for which the basis function expansion implied by the kernel isn't even finite dimensional!

# Examples of Kernel Functions

- **Degree  $B$  Polynomial Kernel:**  $\mathcal{K}_P(\mathbf{x}, \mathbf{x}') = (\mathbf{x}^T \mathbf{x}' + 1)^B$
- **Gaussian/RBF Kernel:**  $\mathcal{K}_G(\mathbf{x}, \mathbf{x}') = \exp(-\gamma \|\mathbf{x} - \mathbf{x}'\|_2^2)$
- Many more domain-specific kernels for strings, histograms, probability distributions, and other complex structured objects.

# Trade-Offs: Basis Expansion and Kernels

- Linear classifiers are fast and space efficient, but can have high bias.
- Basis expansion requires more space ( $O(NK)$  for data and  $O(K)$  for parameters), but yields non-linear classifiers that have lower bias.
- Kernel SVMs actually require  $O(N^2)$  space for storing all the kernel values during training and have  $O(N)$  parameters. This can still be much lower than  $O(NK)$  for large sets of basis functions. Kernels also yield non-linear classifiers that have lower bias.
- Kernel SVMs often have at least two parameters ( $C$  and a kernel hyperparameter). These need to be set jointly, which can be computationally expensive.

# Trade-Offs: Basis Expansion and Kernels

- Gaussian kernel SVMs also have infinite capacity, and a very closely related to weighted KNN.
- Importantly, everything we said about SVMs and kernels is also true for logistic regression. Applying the kernel trick to logistic regression yields a model called “kernel logistic regression” or KLR.
- KLR can exploit infinite dimensional feature spaces, can be learned with smooth optimization methods, supports probabilistic outputs, and has an easy multi-class generalization, but lacks the margin maximization property.