

COMPSCI 589

Lecture 3: Naive Bayes, LDA, and Logistic Regression

Benjamin M. Marlin

College of Information and Computer Sciences
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).

Created with support from National Science Foundation Award# IIS-1350522.

Outline

- 1 Review
- 2 Bayes Optimal Classifiers
- 3 Naive Bayes
- 4 LDA
- 5 Logistic Regression

The Classification Task

Definition: The Classification Task

Given a feature vector $\mathbf{x} \in \mathbb{R}^D$ that describes an object that belongs to one of C classes from the set \mathcal{Y} , predict which class the object belongs to.

The Classifier Learning Problem

Definition: Classifier Learning

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ where $\mathbf{x}_i \in \mathbb{R}^D$ is a feature vector and $y_i \in \mathcal{Y}$ is a class label, learn a function $f : \mathbb{R}^D \rightarrow \mathcal{Y}$ that accurately predicts the class label y for any feature vector \mathbf{x} .

Classification Error and Accuracy

Definition: Classification Error Rate

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ and a function $f : \mathbb{R}^D \rightarrow \mathcal{Y}$, the classification error rate of f on \mathcal{D} is:

$$Err(f, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i \neq f(\mathbf{x}_i))$$

Definition: Classification Accuracy Rate

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ and a function $f : \mathbb{R}^D \rightarrow \mathcal{Y}$, the classification accuracy rate of f on \mathcal{D} is:

$$Acc(f, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N \mathbb{I}(y_i = f(\mathbf{x}_i))$$

Outline

1 Review

2 Bayes Optimal Classifiers

3 Naive Bayes

4 LDA

5 Logistic Regression

Probabilistic Classification

- Suppose we know that the true probability of seeing a data case that belongs to class c is $P(Y = c) = \pi_c$ and the true probability density of seeing a data vector $\mathbf{x} \in \mathbb{R}^D$ that belongs to class c is $p(\mathbf{X} = \mathbf{x} | Y = c) = \phi_c(\mathbf{x})$.

Probabilistic Classification

- Suppose we know that the true probability of seeing a data case that belongs to class c is $P(Y = c) = \pi_c$ and the true probability density of seeing a data vector $\mathbf{x} \in \mathbb{R}^D$ that belongs to class c is $p(\mathbf{X} = \mathbf{x} | Y = c) = \phi_c(\mathbf{x})$.
- We can use this information to compute the probability that the class of any observation \mathbf{x} is c . How?

Probabilistic Classification

- Suppose we know that the true probability of seeing a data case that belongs to class c is $P(Y = c) = \pi_c$ and the true probability density of seeing a data vector $\mathbf{x} \in \mathbb{R}^D$ that belongs to class c is $p(\mathbf{X} = \mathbf{x} | Y = c) = \phi_c(\mathbf{x})$.
- We can use this information to compute the probability that the class of any observation \mathbf{x} is c . How?
- Bayes Rule:
$$P(Y = c | \mathbf{X} = \mathbf{x}) = \frac{\phi_c(\mathbf{x})\pi_c}{\sum_{c' \in \mathcal{Y}} \phi_{c'}(\mathbf{x})\pi_{c'}}$$

The Bayes Optimal Classifier

- The Bayes Optimal Classifier uses the classification function:

$$\begin{aligned}f_B(\mathbf{x}) &= \arg \max_{c \in \mathcal{Y}} P(Y = c | \mathbf{X} = \mathbf{x}) \\&= \arg \max_{c \in \mathcal{Y}} \phi_c(\mathbf{x}) \pi_c\end{aligned}$$

The Bayes Optimal Classifier

- The Bayes Optimal Classifier uses the classification function:

$$\begin{aligned}f_B(\mathbf{x}) &= \arg \max_{c \in \mathcal{Y}} P(Y = c | \mathbf{X} = \mathbf{x}) \\ &= \arg \max_{c \in \mathcal{Y}} \phi_c(\mathbf{x}) \pi_c\end{aligned}$$

- The Bayes optimal classifier achieves the minimal possible expected error rate among all possible classifiers:

$$1 - \mathbb{E}_{P(\mathbf{X}=\mathbf{x})} \left[\max_{c \in \mathcal{Y}} P(Y = c | \mathbf{X} = \mathbf{x}) \right]$$

The Bayes Optimal Classifier

- The Bayes Optimal Classifier uses the classification function:

$$\begin{aligned} f_B(\mathbf{x}) &= \arg \max_{c \in \mathcal{Y}} P(Y = c | \mathbf{X} = \mathbf{x}) \\ &= \arg \max_{c \in \mathcal{Y}} \phi_c(\mathbf{x}) \pi_c \end{aligned}$$

- The Bayes optimal classifier achieves the minimal possible expected error rate among all possible classifiers:

$$1 - \mathbb{E}_{P(\mathbf{X}=\mathbf{x})} \left[\max_{c \in \mathcal{Y}} P(Y = c | \mathbf{X} = \mathbf{x}) \right]$$

- This is a nice result, but is it useful in practice?

Outline

- 1 Review
- 2 Bayes Optimal Classifiers
- 3 Naive Bayes**
- 4 LDA
- 5 Logistic Regression

Naive Bayes

- The naive Bayes classifier approximates the Bayes optimal classifier using a simple form for the functions $\phi_c(\mathbf{x})$.

Naive Bayes

- The naive Bayes classifier approximates the Bayes optimal classifier using a simple form for the functions $\phi_c(\mathbf{x})$.
- This form assumes that all of the data dimensions are probabilistically independent given the value of the class variable:

$$\phi_c(\mathbf{x}) = p(\mathbf{X} = \mathbf{x} | Y = c) = \prod_{d=1}^D p(X_d = x_d | Y = c) = \prod_{d=1}^D \phi_{cd}(x_d)$$

Naive Bayes

- The naive Bayes classifier approximates the Bayes optimal classifier using a simple form for the functions $\phi_c(\mathbf{x})$.
- This form assumes that all of the data dimensions are probabilistically independent given the value of the class variable:

$$\phi_c(\mathbf{x}) = p(\mathbf{X} = \mathbf{x} | Y = c) = \prod_{d=1}^D p(X_d = x_d | Y = c) = \prod_{d=1}^D \phi_{cd}(x_d)$$

- The general form for the classification function is:

$$f_{NB}(\mathbf{x}) = \arg \max_{c \in \mathcal{Y}} \pi_c \prod_{d=1}^D \phi_{cd}(x_d)$$

Class Conditional Distributions

- The functions $p(X_d = x_d | Y = c) = \phi_{cd}(x_d)$ are called *marginal class conditional distributions*.

Class Conditional Distributions

- The functions $p(X_d = x_d | Y = c) = \phi_{cd}(x_d)$ are called *marginal class conditional distributions*.
- For real valued x_d , $p(X_d = x_d | Y = c)$ is typically modeled as a normal density $\phi_{cd}(x_d) = \mathcal{N}(x_d; \mu_{dc}, \sigma_{dc}^2) = \frac{1}{\sqrt{2\pi\sigma_{dc}^2}} \exp\left(-\frac{1}{2\sigma_{dc}^2}(x_d - \mu_{dc})^2\right)$.

Class Conditional Distributions

- The functions $p(X_d = x_d | Y = c) = \phi_{cd}(x_d)$ are called *marginal class conditional distributions*.
- For real valued x_d , $p(X_d = x_d | Y = c)$ is typically modeled as a normal density $\phi_{cd}(x_d) = \mathcal{N}(x_d; \mu_{dc}, \sigma_{dc}^2) = \frac{1}{\sqrt{2\pi\sigma_{dc}^2}} \exp\left(-\frac{1}{2\sigma_{dc}^2}(x_d - \mu_{dc})^2\right)$.
- For binary valued x_d , $p(X_d = x_d | Y = c)$ is a Bernoulli distribution $\phi_{cd}(x_d) = \theta_{dc}^{x_d} (1 - \theta_{dc})^{(1-x_d)}$.

Class Conditional Distributions

- The functions $p(X_d = x_d | Y = c) = \phi_{cd}(x_d)$ are called *marginal class conditional distributions*.
- For real valued x_d , $p(X_d = x_d | Y = c)$ is typically modeled as a normal density $\phi_{cd}(x_d) = \mathcal{N}(x_d; \mu_{dc}, \sigma_{dc}^2) = \frac{1}{\sqrt{2\pi\sigma_{dc}^2}} \exp\left(-\frac{1}{2\sigma_{dc}^2}(x_d - \mu_{dc})^2\right)$.
- For binary valued x_d , $p(X_d = x_d | Y = c)$ is a Bernoulli distribution $\phi_{cd}(x_d) = \theta_{dc}^{x_d} (1 - \theta_{dc})^{(1-x_d)}$.
- For general categorical values x_d , $p(X_d = x_d | Y = c)$ is a categorical distribution $\phi_{cd}(x_d) = \prod_{v \in \mathcal{X}_d} \theta_{vdc}^{[x_d=v]}$

Learning for Naive Bayes

- The class probabilities π_c and the parameters of the marginal class conditional distributions $\phi_{cd}(x_d)$ are learned by maximum likelihood on a sample of training data $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$.

Learning for Naive Bayes

- The class probabilities π_c and the parameters of the marginal class conditional distributions $\phi_{cd}(x_d)$ are learned by maximum likelihood on a sample of training data $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$.
- This reduces to estimating them using class-conditional sample averages for most common distributions.

Learning for Naive Bayes

- The class probabilities π_c and the parameters of the marginal class conditional distributions $\phi_{cd}(x_d)$ are learned by maximum likelihood on a sample of training data $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$.
- This reduces to estimating them using class-conditional sample averages for most common distributions.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

Learning for Naive Bayes

- The class probabilities π_c and the parameters of the marginal class conditional distributions $\phi_{cd}(x_d)$ are learned by maximum likelihood on a sample of training data $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$.
- This reduces to estimating them using class-conditional sample averages for most common distributions.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

- Normal: $\mu_{dc} = \frac{\sum_{i=1}^n [y_i=c] x_{di}}{\sum_{i=1}^n [y_i=c]}$ $\sigma_{dc}^2 = \frac{\sum_{i=1}^n [y_i=c] (x_{di} - \mu_{dc})^2}{\sum_{i=1}^n [y_i=c]}$

Learning for Naive Bayes

- The class probabilities π_c and the parameters of the marginal class conditional distributions $\phi_{cd}(x_d)$ are learned by maximum likelihood on a sample of training data $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$.
- This reduces to estimating them using class-conditional sample averages for most common distributions.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

- Normal: $\mu_{dc} = \frac{\sum_{i=1}^n [y_i=c] x_{di}}{\sum_{i=1}^n [y_i=c]}$ $\sigma_{dc}^2 = \frac{\sum_{i=1}^n [y_i=c] (x_{di} - \mu_{dc})^2}{\sum_{i=1}^n [y_i=c]}$

- Bernoulli: $\theta_{dc} = \frac{\sum_{i=1}^n [y_i=c] x_{di}}{\sum_{i=1}^n [y_i=c]}$ Categorical: $\theta_{vdc} = \frac{\sum_{i=1}^n [y_i=c] [x_{di}=v]}{\sum_{i=1}^n [y_i=c]}$

Geometric Interpretation

- Suppose we use normal distributions for the marginal class conditional distributions and we have a binary classification problem. What is the geometry of the decision boundary?

Geometric Interpretation

- Suppose we use normal distributions for the marginal class conditional distributions and we have a binary classification problem. What is the geometry of the decision boundary?

$$\begin{aligned}f_{NB}(\mathbf{x}) &= \arg \max_{c \in \mathcal{Y}} \pi_c \prod_{d=1}^D \phi_{cd}(x_d) \\&= \arg \max_{c \in \mathcal{Y}} \log(\pi_c) + \sum_{d=1}^D \log(\phi_{cd}(x_d)) \\&= \arg \max_{c \in \mathcal{Y}} \log(\pi_c) + \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{dc}^2) - \frac{1}{2\sigma_{dc}^2} (x_d - \mu_{dc})^2 \right)\end{aligned}$$

Geometric Interpretation

- The decision boundary consists of the set of points \mathbf{x} where:

Geometric Interpretation

- The decision boundary consists of the set of points \mathbf{x} where:

$$\log(\pi_0) + \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d0}^2) - \frac{1}{2\sigma_{d0}^2} (x_d - \mu_{d0})^2 \right) \\ - \log(\pi_1) - \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d1}^2) - \frac{1}{2\sigma_{d1}^2} (x_d - \mu_{d1})^2 \right) = 0$$

Geometric Interpretation

- The decision boundary consists of the set of points \mathbf{x} where:

$$\begin{aligned} \log(\pi_0) + \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d0}^2) - \frac{1}{2\sigma_{d0}^2} (x_d - \mu_{d0})^2 \right) \\ - \log(\pi_1) - \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d1}^2) - \frac{1}{2\sigma_{d1}^2} (x_d - \mu_{d1})^2 \right) = 0 \end{aligned}$$

- It's easy to see that the decision boundary is a quadratic function of \mathbf{x} with the form: $\sum_{d=1}^D (a_d x_d^2 + b_d x_d) + c = 0$.

Geometric Interpretation

- The decision boundary consists of the set of points \mathbf{x} where:

$$\begin{aligned} \log(\pi_0) + \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d0}^2) - \frac{1}{2\sigma_{d0}^2} (x_d - \mu_{d0})^2 \right) \\ - \log(\pi_1) - \sum_{d=1}^D \left(-\frac{1}{2} \log(2\pi\sigma_{d1}^2) - \frac{1}{2\sigma_{d1}^2} (x_d - \mu_{d1})^2 \right) = 0 \end{aligned}$$

- It's easy to see that the decision boundary is a quadratic function of \mathbf{x} with the form: $\sum_{d=1}^D (a_d x_d^2 + b_d x_d) + c = 0$.
- In the multi-class case, the decision boundary is piece-wise quadratic.

Trade-Offs

- Speed: Both learning and classification have very low computational complexity.

Trade-Offs

- Speed: Both learning and classification have very low computational complexity.
- Storage: The model only requires $O(DC)$ parameters. It achieves a large compression of the training data.

Trade-Offs

- Speed: Both learning and classification have very low computational complexity.
- Storage: The model only requires $O(DC)$ parameters. It achieves a large compression of the training data.
- Interpretability: The model has good interpretability since the parameters of $\phi_c(\mathbf{x})$ correspond to class conditional averages.

Trade-Offs

- Speed: Both learning and classification have very low computational complexity.
- Storage: The model only requires $O(DC)$ parameters. It achieves a large compression of the training data.
- Interpretability: The model has good interpretability since the parameters of $\phi_c(\mathbf{x})$ correspond to class conditional averages.
- Accuracy: The assumption of feature independence and canonical forms for class-conditional marginals will rarely be correct for real-world problems, leading to lower accuracy.

Trade-Offs

- Speed: Both learning and classification have very low computational complexity.
- Storage: The model only requires $O(DC)$ parameters. It achieves a large compression of the training data.
- Interpretability: The model has good interpretability since the parameters of $\phi_c(\mathbf{x})$ correspond to class conditional averages.
- Accuracy: The assumption of feature independence and canonical forms for class-conditional marginals will rarely be correct for real-world problems, leading to lower accuracy.
- Data: Some care is needed in the estimation of parameters in the discrete case when data is scarce.

Outline

- 1 Review
- 2 Bayes Optimal Classifiers
- 3 Naive Bayes
- 4 LDA**
- 5 Logistic Regression

Linear Discriminant Analysis

- Linear Discriminant Analysis (LDA) is a classification technique due to Fisher that dates back to the 1930's.

Linear Discriminant Analysis

- Linear Discriminant Analysis (LDA) is a classification technique due to Fisher that dates back to the 1930's.
- It can be interpreted as a different approximation to the Bayes Optimal Classifier for real-valued data.

Linear Discriminant Analysis

- Linear Discriminant Analysis (LDA) is a classification technique due to Fisher that dates back to the 1930's.
- It can be interpreted as a different approximation to the Bayes Optimal Classifier for real-valued data.
- Instead of a product of independent Normals like in Naive Bayes, LDA assumes the class-conditional densities are multivariate normal with a common covariance matrix:

$$\phi_c(\mathbf{x}) = \mathcal{N}(\mathbf{x}; \mu_c, \Sigma) = \frac{1}{|2\pi\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu_c)^T \Sigma^{-1}(\mathbf{x} - \mu_c)\right)$$

Linear Discriminant Analysis

- Linear Discriminant Analysis (LDA) is a classification technique due to Fisher that dates back to the 1930's.
- It can be interpreted as a different approximation to the Bayes Optimal Classifier for real-valued data.
- Instead of a product of independent Normals like in Naive Bayes, LDA assumes the class-conditional densities are multivariate normal with a common covariance matrix:

$$\phi_c(\mathbf{x}) = \mathcal{N}(\mathbf{x}; \mu_c, \Sigma) = \frac{1}{|2\pi\Sigma|^{1/2}} \exp\left(-\frac{1}{2}(\mathbf{x} - \mu_c)^T \Sigma^{-1}(\mathbf{x} - \mu_c)\right)$$

- The classification function is $f_{LDA}(\mathbf{x}) = \arg \max_{c \in \mathcal{Y}} \phi_c(\mathbf{x}) \pi_c$

Learning for LDA

- As with Naive Bayes, LDA parameters are learned using maximum likelihood, which reduces to using sample estimates.

Learning for LDA

- As with Naive Bayes, LDA parameters are learned using maximum likelihood, which reduces to using sample estimates.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

Learning for LDA

- As with Naive Bayes, LDA parameters are learned using maximum likelihood, which reduces to using sample estimates.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

- Class Means: $\mu_c = \frac{\sum_{i=1}^n [y_i=c] \mathbf{x}_i}{\sum_{i=1}^n [y_i=c]}$

Learning for LDA

- As with Naive Bayes, LDA parameters are learned using maximum likelihood, which reduces to using sample estimates.

- Class probabilities: $\pi_c = \frac{1}{n} \sum_{i=1}^n [y_i = c]$

- Class Means: $\mu_c = \frac{\sum_{i=1}^n [y_i=c] \mathbf{x}_i}{\sum_{i=1}^n [y_i=c]}$

- Shared Covariance: $\Sigma = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \mu_{y_i})(\mathbf{x}_i - \mu_{y_i})^T$

Geometric Interpretation

- What is the geometry of the decision boundary for LDA in the binary case?

Geometric Interpretation

- What is the geometry of the decision boundary for LDA in the binary case?
- The decision boundary consists of the set of points \mathbf{x} where:

$$\begin{aligned} \log(\pi_0) - \frac{1}{2} \log |2\pi\Sigma| - \frac{1}{2}(\mathbf{x} - \mu_0)^T \Sigma^{-1}(\mathbf{x} - \mu_0) \\ - \log(\pi_1) + \frac{1}{2} \log |2\pi\Sigma| + \frac{1}{2}(\mathbf{x} - \mu_1)^T \Sigma^{-1}(\mathbf{x} - \mu_1) = 0 \end{aligned}$$

Geometric Interpretation

- What is the geometry of the decision boundary for LDA in the binary case?
- The decision boundary consists of the set of points \mathbf{x} where:

$$\begin{aligned} \log(\pi_0) - \frac{1}{2} \log |2\pi\Sigma| - \frac{1}{2}(\mathbf{x} - \mu_0)^T \Sigma^{-1}(\mathbf{x} - \mu_0) \\ - \log(\pi_1) + \frac{1}{2} \log |2\pi\Sigma| + \frac{1}{2}(\mathbf{x} - \mu_1)^T \Sigma^{-1}(\mathbf{x} - \mu_1) = 0 \end{aligned}$$

- We can cancel a large number of terms because of the common covariance matrix and obtain the following result:

$$\log(\pi_0) - \log(\pi_1) - 0.5\mu_0^T \Sigma^{-1} \mu_0 + 0.5\mu_1^T \Sigma^{-1} \mu_1 + (\mu_0 - \mu_1)^T \Sigma^{-1} \mathbf{x} = 0$$

Geometric Interpretation

- What is the geometry of the decision boundary for LDA in the binary case?
- The decision boundary consists of the set of points \mathbf{x} where:

$$\begin{aligned} \log(\pi_0) - \frac{1}{2} \log |2\pi\Sigma| - \frac{1}{2}(\mathbf{x} - \mu_0)^T \Sigma^{-1}(\mathbf{x} - \mu_0) \\ - \log(\pi_1) + \frac{1}{2} \log |2\pi\Sigma| + \frac{1}{2}(\mathbf{x} - \mu_1)^T \Sigma^{-1}(\mathbf{x} - \mu_1) = 0 \end{aligned}$$

- We can cancel a large number of terms because of the common covariance matrix and obtain the following result:

$$\log(\pi_0) - \log(\pi_1) - 0.5\mu_0^T \Sigma^{-1} \mu_0 + 0.5\mu_1^T \Sigma^{-1} \mu_1 + (\mu_0 - \mu_1)^T \Sigma^{-1} \mathbf{x} = 0$$

- This shows that the decision boundary is actually linear in \mathbf{x} .

Trade-Offs

- Speed: The quadratic dependence on D makes LDA slower than Naive Bayes by a factor of D during learning and classification.

Trade-Offs

- Speed: The quadratic dependence on D makes LDA slower than Naive Bayes by a factor of D during learning and classification.
- Storage: The model requires $O(D^2C)$ parameters. This can still represent a good compression of the data when $D \ll N$.

Trade-Offs

- Speed: The quadratic dependence on D makes LDA slower than Naive Bayes by a factor of D during learning and classification.
- Storage: The model requires $O(D^2C)$ parameters. This can still represent a good compression of the data when $D \ll N$.
- Interpretability: The model has good interpretability since the mean parameters μ_c correspond to class conditional averages.

Trade-Offs

- Speed: The quadratic dependence on D makes LDA slower than Naive Bayes by a factor of D during learning and classification.
- Storage: The model requires $O(D^2C)$ parameters. This can still represent a good compression of the data when $D \ll N$.
- Interpretability: The model has good interpretability since the mean parameters μ_c correspond to class conditional averages.
- Accuracy: The assumptions LDA makes will rarely be correct for real-world problems. However, the induced linear decision boundaries can often perform reasonably well.

Trade-Offs

- Speed: The quadratic dependence on D makes LDA slower than Naive Bayes by a factor of D during learning and classification.
- Storage: The model requires $O(D^2C)$ parameters. This can still represent a good compression of the data when $D \ll N$.
- Interpretability: The model has good interpretability since the mean parameters μ_c correspond to class conditional averages.
- Accuracy: The assumptions LDA makes will rarely be correct for real-world problems. However, the induced linear decision boundaries can often perform reasonably well.
- Data: LDA will generally need more data than NB since it needs to estimate the $O(D^2)$ parameters in the pooled covariance matrix Σ .

Outline

- 1 Review
- 2 Bayes Optimal Classifiers
- 3 Naive Bayes
- 4 LDA
- 5 Logistic Regression**

Generative vs Discriminative Classifiers

- The Bayes Optimal Classifier, Naive Bayes and LDA are said to be *generative* classifiers because they explicitly model the joint distribution $P(\mathbf{X}, Y)$ of the data vectors \mathbf{x} and the labels y .

Generative vs Discriminative Classifiers

- The Bayes Optimal Classifier, Naive Bayes and LDA are said to be *generative* classifiers because they explicitly model the joint distribution $P(\mathbf{X}, Y)$ of the data vectors \mathbf{x} and the labels y .
Question: is this really necessary?

Generative vs Discriminative Classifiers

- The Bayes Optimal Classifier, Naive Bayes and LDA are said to be *generative* classifiers because they explicitly model the joint distribution $P(\mathbf{X}, Y)$ of the data vectors \mathbf{x} and the labels y .
Question: is this really necessary?
- No, to build a probabilistic classifier, all we really need to model is $P(Y|\mathbf{X})$.

Generative vs Discriminative Classifiers

- The Bayes Optimal Classifier, Naive Bayes and LDA are said to be *generative* classifiers because they explicitly model the joint distribution $P(\mathbf{X}, Y)$ of the data vectors \mathbf{x} and the labels y .
Question: is this really necessary?
- No, to build a probabilistic classifier, all we really need to model is $P(Y|\mathbf{X})$.
- Classifiers based on directly estimating $P(Y|\mathbf{X})$ are called *discriminative* classifiers because they ignore the distribution of \mathbf{x} and focus only on the class labels y .

Logistic Regression

- Logistic Regression is a probabilistic discriminative classifier.

Logistic Regression

- Logistic Regression is a probabilistic discriminative classifier.
- In the binary case, it directly models the decision boundary using a linear function:

Logistic Regression

- Logistic Regression is a probabilistic discriminative classifier.
- In the binary case, it directly models the decision boundary using a linear function:

$$\log P(Y = 1|\mathbf{x}) - \log P(Y = 0|\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

Logistic Regression

- Logistic Regression is a probabilistic discriminative classifier.
- In the binary case, it directly models the decision boundary using a linear function:

$$\log P(Y = 1|\mathbf{x}) - \log P(Y = 0|\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

$$P(Y = 1|\mathbf{X}) = \frac{1}{1 + \exp(-(\mathbf{w}^T \mathbf{x} + b))}$$

Logistic Regression


- Logistic Regression is a probabilistic discriminative classifier.
- In the binary case, it directly models the decision boundary using a linear function:

$$\log P(Y = 1|\mathbf{x}) - \log P(Y = 0|\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

$$P(Y = 1|\mathbf{X}) = \frac{1}{1 + \exp(-(\mathbf{w}^T \mathbf{x} + b))}$$

- The classification function is: $f_{LR}(\mathbf{x}) = \arg \max_{c \in \mathcal{Y}} P(Y = c|\mathbf{x})$

Logistic Function



Multiclass Logistic Regression

- Logistic regression can also be extended to the multiclass case:

$$P(Y = c|\mathbf{x}) = \frac{\exp(\mathbf{w}_c^T \mathbf{x} + b_c)}{\sum_{c' \in \mathcal{Y}} \exp(-(\mathbf{w}_{c'}^T \mathbf{x} + b_{c'}))}$$

Multiclass Logistic Regression

- Logistic regression can also be extended to the multiclass case:

$$P(Y = c|\mathbf{x}) = \frac{\exp(\mathbf{w}_c^T \mathbf{x} + b_c)}{\sum_{c' \in \mathcal{Y}} \exp(\mathbf{w}_{c'}^T \mathbf{x} + b_{c'})}$$

- The classification function is still:

$$f_{LR}(\mathbf{x}) = \arg \max_{c \in \mathcal{Y}} P(Y = c|\mathbf{x})$$

Learning Logistic Regression

- The logistic regression model parameters $\theta = \{(\mathbf{w}_c, b_c), c \in \mathcal{Y}\}$ are selected to optimize the conditional log likelihood of the labels given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$:

Learning Logistic Regression

- The logistic regression model parameters $\theta = \{(\mathbf{w}_c, b_c), c \in \mathcal{Y}\}$ are selected to optimize the conditional log likelihood of the labels given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$:

$$\theta_* = \arg \max_{\theta} \mathcal{L}(\theta | \mathcal{D}) = \arg \max_{\theta} \sum_{i=1}^n \log P(Y = y_i | \mathbf{X} = \mathbf{x}_i)$$

Learning Logistic Regression

- The logistic regression model parameters $\theta = \{(\mathbf{w}_c, b_c), c \in \mathcal{Y}\}$ are selected to optimize the conditional log likelihood of the labels given a data set $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : n\}$:

$$\theta_* = \arg \max_{\theta} \mathcal{L}(\theta | \mathcal{D}) = \arg \max_{\theta} \sum_{i=1}^n \log P(Y = y_i | \mathbf{X} = \mathbf{x}_i)$$

- However, the function $\mathcal{L}(\theta | \mathcal{D})$ can not be maximized analytically. Learning the model parameters requires numerical optimization methods.

Geometry

- Logistic regression is explicitly designed to have a linear decision boundary in the binary case.

Geometry

- Logistic regression is explicitly designed to have a linear decision boundary in the binary case.
- In the multiclass case, the decision boundary is piece-wise linear.

Geometry

- Logistic regression is explicitly designed to have a linear decision boundary in the binary case.
- In the multiclass case, the decision boundary is piece-wise linear.
- Logistic Regression has the same representational capacity as LDA.

Trade-Offs

- Speed: At classification time, LR is faster than NB and LDA. Learning LR requires numerical optimization, which will be slower than NB.

Trade-Offs

- Speed: At classification time, LR is faster than NB and LDA. Learning LR requires numerical optimization, which will be slower than NB.
- Storage: The model requires $O(DC)$ parameters. The same order as Naive Bayes, but much less than LDA's $O(DC + D^2)$ when $C \ll D$.

Trade-Offs

- Speed: At classification time, LR is faster than NB and LDA. Learning LR requires numerical optimization, which will be slower than NB.
- Storage: The model requires $O(DC)$ parameters. The same order as Naive Bayes, but much less than LDA's $O(DC + D^2)$ when $C \ll D$.
- Interpretability: The “importance” of different feature variables x_d can be understood in terms of their weights w_{dc} .

Trade-Offs

- Speed: At classification time, LR is faster than NB and LDA. Learning LR requires numerical optimization, which will be slower than NB.
- Storage: The model requires $O(DC)$ parameters. The same order as Naive Bayes, but much less than LDA's $O(DC + D^2)$ when $C \ll D$.
- Interpretability: The “importance” of different feature variables x_d can be understood in terms of their weights w_{dc} .
- Accuracy: Tends to be better in high dimensions with limited data compared to LDA. Much worse than KNN in low dimensions with lots of data and non-linear decision boundaries.