

COMPSCI 589

Lecture 9: KNN Regression, Regression Trees, and Feature Selection

Benjamin M. Marlin

College of Information and Computer Sciences
University of Massachusetts Amherst

Slides by Benjamin M. Marlin (marlin@cs.umass.edu).
Created with support from National Science Foundation Award# IIS-1350522.


Outline

- 1 Review
- 2 KNN Regression
- 3 Regression Trees
- 4 Feature Selection

The Regression Task

Definition: The Regression Task

Given a feature vector $\mathbf{x} \in \mathbb{R}^D$, predict it's corresponding output value y .



The Regression Learning Problem

Definition: Regression Learning Problem

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ where $\mathbf{x}_i \in \mathbb{R}^D$ is a feature vector and $y_i \in \mathbb{R}$ is the output, learn a function $f : \mathbb{R}^D \rightarrow \mathbb{R}$ that accurately predicts y for any feature vector \mathbf{x} .

Error Measures: MSE

Definition: Mean Squared Error

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ and a function $f : \mathbb{R}^D \rightarrow \mathcal{Y}$, the mean squared error of f on \mathcal{D} is:

$$MSE(f, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2$$

Error Measures: MSE

Definition: Mean Squared Error

Given a data set of example pairs $\mathcal{D} = \{(\mathbf{x}_i, y_i), i = 1 : N\}$ and a function $f : \mathbb{R}^D \rightarrow \mathcal{Y}$, the mean squared error of f on \mathcal{D} is:

$$MSE(f, \mathcal{D}) = \frac{1}{N} \sum_{i=1}^N (y_i - f(\mathbf{x}_i))^2$$

Related measures include:

Sum of Squared Errors: $SSE(f, \mathcal{D}) = N \cdot MSE(f, \mathcal{D})$

Risidual Sum of Squares: $RSS(f, \mathcal{D}) = N \cdot MSE(f, \mathcal{D})$

Root Mean Squared Error: $RMSE(f, \mathcal{D}) = \sqrt{MSE(f, \mathcal{D})}$

Outline

- 1 Review
- 2 KNN Regression
- 3 Regression Trees
- 4 Feature Selection

K Nearest Neighbors Regression

The KNN regression is a non-parametric regression method that simply stores the training data \mathcal{D} and makes a prediction for each new instance \mathbf{x} using an average over it's set of K nearest neighbors $\mathcal{N}_K(\mathbf{x})$ computed using any distance function $d : \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$.

K Nearest Neighbors Regression

The KNN regression is a non-parametric regression method that simply stores the training data \mathcal{D} and makes a prediction for each new instance \mathbf{x} using an average over it's set of K nearest neighbors $\mathcal{N}_K(\mathbf{x})$ computed using any distance function $d : \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$.

KNN Regression Function

$$f_{KNN}(\mathbf{x}) = \frac{1}{K} \sum_{i \in \mathcal{N}_K(\mathbf{x})} y_i$$

K Nearest Neighbors Regression



The KNN regression is a non-parametric regression method that simply stores the training data \mathcal{D} and makes a prediction for each new instance \mathbf{x} using an average over it's set of K nearest neighbors $\mathcal{N}_K(\mathbf{x})$ computed using any distance function $d : \mathbb{R}^D \times \mathbb{R}^D \rightarrow \mathbb{R}$.

KNN Regression Function



$$f_{KNN}(\mathbf{x}) = \frac{1}{K} \sum_{i \in \mathcal{N}_K(\mathbf{x})} y_i$$

As with classification, use of KNN requires choosing the distance function d and the number of neighbors K .

Example: 1D KNN (K=1 vs K=9)



Example: 2D KNN (K=1 vs K=9)



Weighted KNN Regression

- Instead of giving all of the K neighbors equal weight in the average, a distance-weighted average can be used:

$$f_{KNN}(\mathbf{x}) = \frac{\sum_{i \in \mathcal{N}_K(\mathbf{x})} w_i y_i}{\sum_{i \in \mathcal{N}_K(\mathbf{x})} w_i}$$
$$w_i = \exp(-\alpha d_i)$$

Outline

- 1 Review
- 2 KNN Regression
- 3 Regression Trees**
- 4 Feature Selection

Regression Trees

- A regression tree makes predictions using a conjunction of rules organized into a binary tree structure.



Regression Trees

- A regression tree makes predictions using a conjunction of rules organized into a binary tree structure.
- Each internal node in a regression tree contains a rule of the form $(x_d < t)$ or $(x_d = t)$ that tests a single data dimension d against a single threshold value t and assigns the data case to it's left or right sub-tree according to the result.

Regression Trees

- A regression tree makes predictions using a conjunction of rules organized into a binary tree structure.
- Each internal node in a regression tree contains a rule of the form $(x_d < t)$ or $(x_d = t)$ that tests a single data dimension d against a single threshold value t and assigns the data case to its left or right sub-tree according to the result.
- A data case is routed through the tree from the root to a leaf. Each leaf node is associated with a predicted output, and a data case is assigned the output of the leaf node it is routed to.

Example: 2D Regression Trees



Building Regression Trees

Algorithm 1 $BuildTree(Root, \mathcal{D}, h, minS, maxD)$

$d, t = BestSplit(\mathcal{D})$

Building Regression Trees

Algorithm 2 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

 $d, t = \text{BestSplit}(\mathcal{D})$ $\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

Building Regression Trees

Algorithm 3 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

 $d, t = \text{BestSplit}(\mathcal{D})$ $\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$ **if** $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

Building Regression Trees

Algorithm 4 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

 $d, t = \text{BestSplit}(\mathcal{D})$ $\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$ **if** $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then** $\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

Building Regression Trees

Algorithm 5 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

Building Regression Trees

Algorithm 6 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}$, $\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

Building Regression Trees

Algorithm 7 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

Building Regression Trees

Algorithm 8 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

$\text{BuildTree}(\text{Root.RightChild}, \mathcal{D}_1, h + 1, \text{minS}, \text{maxD})$

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

Building Regression Trees

Algorithm 9 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

$\text{BuildTree}(\text{Root.RightChild}, \mathcal{D}_1, h + 1, \text{minS}, \text{maxD})$

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

Building Regression Trees

Algorithm 10 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

else

Building Regression Trees

Algorithm 11 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

else

BuildTree(*Root.LeftChild*, \mathcal{D}_2 , $h + 1$, *minS*, *maxD*)

Building Regression Trees

Algorithm 12 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}$, $\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

else

BuildTree(*Root.LeftChild*, \mathcal{D}_2 , $h + 1$, *minS*, *maxD*)

Building Regression Trees

Algorithm 13 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

else

BuildTree(*Root.LeftChild*, \mathcal{D}_2 , $h + 1$, *minS*, *maxD*)

$\text{Root.d} = d, \text{Root.t} = t$

Building Regression Trees

Algorithm 14 *BuildTree*(*Root*, \mathcal{D} , *h*, *minS*, *maxD*)

$d, t = \text{BestSplit}(\mathcal{D})$

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}, \mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

if $|\mathcal{D}_1| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.RightChild.Prediction} = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

else

BuildTree(*Root.RightChild*, \mathcal{D}_1 , $h + 1$, *minS*, *maxD*)

if $|\mathcal{D}_2| \leq \text{minS}$ or $h + 1 \geq \text{maxD}$ **then**

$\text{Root.LeftChild.Prediction} = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

else

BuildTree(*Root.LeftChild*, \mathcal{D}_2 , $h + 1$, *minS*, *maxD*)

$\text{Root.d} = d, \text{Root.t} = t$

return *Root*

Finding the Best Split

Algorithm 15 $BestSplit(\mathcal{D})$

Finding the Best Split

Algorithm 16 $BestSplit(\mathcal{D})$

for d from 1 to D **do**

Finding the Best Split

Algorithm 17 *BestSplit*(\mathcal{D})

for d from 1 to D **do** $\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

Finding the Best Split

Algorithm 18 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

Finding the Best Split

Algorithm 19 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

Finding the Best Split

Algorithm 20 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

Finding the Best Split

Algorithm 21 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

Finding the Best Split

Algorithm 22 $BestSplit(\mathcal{D})$

for d from 1 to D **do**

$\mathbf{s} = sort(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 | i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) | x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) | x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

Finding the Best Split

Algorithm 23 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

$\text{Score}(d, t) = \sum_{y \in \mathcal{D}_1} (y - \bar{y}_1)^2 + \sum_{y \in \mathcal{D}_2} (y - \bar{y}_2)^2$

Finding the Best Split

Algorithm 24 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

$\text{Score}(d, t) = \sum_{y \in \mathcal{D}_1} (y - \bar{y}_1)^2 + \sum_{y \in \mathcal{D}_2} (y - \bar{y}_2)^2$

Finding the Best Split

Algorithm 25 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

$\text{Score}(d, t) = \sum_{y \in \mathcal{D}_1} (y - \bar{y}_1)^2 + \sum_{y \in \mathcal{D}_2} (y - \bar{y}_2)^2$

Finding the Best Split

Algorithm 26 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

$\text{Score}(d, t) = \sum_{y \in \mathcal{D}_1} (y - \bar{y}_1)^2 + \sum_{y \in \mathcal{D}_2} (y - \bar{y}_2)^2$

$d, t = \arg \min_{d', t'} \text{Score}(d', t')$

Finding the Best Split

Algorithm 27 *BestSplit*(\mathcal{D})

for d from 1 to D **do**

$\mathbf{s} = \text{sort}(\{x_{d1}, \dots, x_{dN}\})$

for t in $\{(s_i + s_{i+1})/2 \mid i = 1 \dots N - 1\}$ **do**

$\mathcal{D}_1 = \{(y_i, \mathbf{x}_i) \mid x_{di} \leq t\}$

$\mathcal{D}_2 = \{(y_i, \mathbf{x}_i) \mid x_{di} > t\}$

$\bar{y}_1 = \frac{1}{|\mathcal{D}_1|} \sum_{y \in \mathcal{D}_1} y$

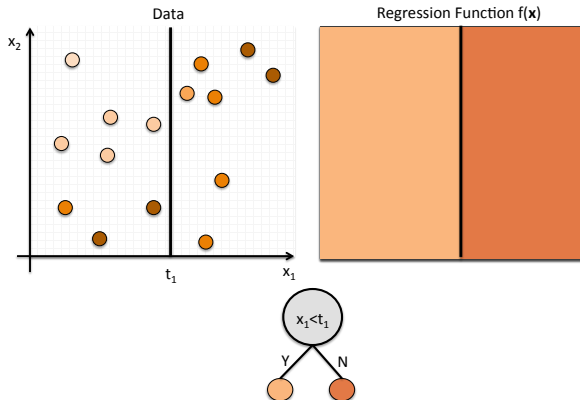
$\bar{y}_2 = \frac{1}{|\mathcal{D}_2|} \sum_{y \in \mathcal{D}_2} y$

$\text{Score}(d, t) = \sum_{y \in \mathcal{D}_1} (y - \bar{y}_1)^2 + \sum_{y \in \mathcal{D}_2} (y - \bar{y}_2)^2$

$d, t = \arg \min_{d', t'} \text{Score}(d', t')$

return (d, t)

Example: Building Regression Trees



Outline

- 1 Review
- 2 KNN Regression
- 3 Regression Trees
- 4 Feature Selection**

Best Subset Selection

Algorithm 6.1 *Best subset selection*

1. Let \mathcal{M}_0 denote the *null model*, which contains no predictors. This model simply predicts the sample mean for each observation.
 2. For $k = 1, 2, \dots, p$:
 - (a) Fit all $\binom{p}{k}$ models that contain exactly k predictors.
 - (b) Pick the best among these $\binom{p}{k}$ models, and call it \mathcal{M}_k . Here *best* is defined as having the smallest RSS, or equivalently largest R^2 .
 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error, C_p (AIC), BIC, or adjusted R^2 .
-

Forward Stepwise Selection

Algorithm 6.2 *Forward stepwise selection*

1. Let \mathcal{M}_0 denote the *null* model, which contains no predictors.
 2. For $k = 0, \dots, p - 1$:
 - (a) Consider all $p - k$ models that augment the predictors in \mathcal{M}_k with one additional predictor.
 - (b) Choose the *best* among these $p - k$ models, and call it \mathcal{M}_{k+1} . Here *best* is defined as having smallest RSS or highest R^2 .
 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error, C_p (AIC), BIC, or adjusted R^2 .
-

Backward Stepwise Selection

Algorithm 6.3 *Backward stepwise selection*

1. Let \mathcal{M}_p denote the *full* model, which contains all p predictors.
 2. For $k = p, p - 1, \dots, 1$:
 - (a) Consider all k models that contain all but one of the predictors in \mathcal{M}_k , for a total of $k - 1$ predictors.
 - (b) Choose the *best* among these k models, and call it \mathcal{M}_{k-1} . Here *best* is defined as having smallest RSS or highest R^2 .
 3. Select a single best model from among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using cross-validated prediction error, C_p (AIC), BIC, or adjusted R^2 .
-